

Version 1.2.1 (19 September 2006)

This document is intended to assist software developers who are writing applications that use TIPC. For information about setting up and operating a network that supports TIPC please see the TIPC User's Guide.

ADVISORY:

Many topics discussed in this document are presented in a very compact format, which presents as much information as possible in as few words as possible. Readers will gain the most benefit from this document by carefully reading (and re-reading) the material, as this will provide a better understanding of TIPC than is likely to be obtained by quickly skimming through it.

Table of Contents

-----

1. TIPC Fundamentals
2. Socket API
3. Native API
4. FAQ
5. Tips and Techniques

1. TIPC Fundamentals

-----

A brief summary of the major concepts is provided in the following sections.

For a comprehensive description of TIPC, please consult the latest version of the TIPC specification at <http://sourceforge.net/projects/tipc>.

IMPORTANT:

At the time of this writing, the TIPC specification has not been updated to include the latest modifications incorporated into TIPC version 1.5 (or later). In cases of conflict between the specification and this document, the information in this document takes precedence.

1.1 TIPC Network Structure

-----

Conceptually, a TIPC network consists of individual processing elements or "nodes". A set of related nodes form a "cluster", while a set of related clusters form a "zone". Typically the grouping of nodes into clusters and zones is based on location; for example, all nodes in the same shelf or the same room may be assigned to the same cluster, while all clusters in the same building may be assigned to the same zone.

Each node in a TIPC network is assigned a unique network address consisting of a zone, cluster, and node identifier, usually denoted <Z.C.N>. Each of these identifiers is an integer value in the range from 1 to the maximum value defined by the network administrator. (Exception: A TIPC node which is not yet part of a larger network uses the default network address of <0.0.0> until a real network address is assigned.) By default the maximum values for zone, cluster, and node identifiers are 3, 1, and 63, respectively, while the theoretical maximums are 255, 4095, and 2047, respectively.

A TIPC node is also assigned a "network identifier". This allows multiple

TIPC networks to use the same physical medium (for example, the same Ethernet cables) without interfering with one another -- each node only recognizes traffic originating on nodes having the same network identifier.

Nodes in a TIPC network communicate with each other by using one or more network interfaces to send and receive messages. Each network interface must be connected to a physical medium that is supported by TIPC (such as Ethernet). When properly configured, TIPC automatically establishes "links" to enable communication with the other nodes in the network, and takes care of routing traffic over the appropriate link, retransmitting messages in the event of errors, etc.

#### THINGS TO REMEMBER:

- TIPC network addressing is NOT like IP network addressing!!! There is only one network address per node in TIPC, even if the node has multiple network interfaces. A node's network interfaces are NOT assigned network addresses at all.
- The network administrator takes care of assigning the network address and the network identifier for each node in the network, so programmers don't have to worry about this.
- The network administrator also takes care of configuring each node's network interfaces to enable communication with all other nodes in the network, so programmers don't have to worry about this either.

#### CURRENT LIMITATIONS:

- TIPC only supports a single cluster per zone.
- TIPC does not support inter-zone communication, so nodes in different zones are effectively part of distinct networks even if they have the same network identifier.
- TIPC does not support the "secondary nodes" concept mentioned in the specification document.

### 1.2 Messaging Overview

---

Applications in a TIPC network typically communicate with one another by sending data units called "messages" between communication endpoints called "ports".

From an application's perspective, a message is a byte string from 1 to 66000 bytes long, whose internal structure is determined by the application. A port is an entity that can send and receive messages in either a connection-oriented manner or a connectionless manner.

Connection-oriented messaging allows a port to establish a connection to a peer port elsewhere in the network, and then exchange messages with that peer. A connection can be established using an explicit handshake mechanism prior to the exchange of any application messages (a variation of the SYN/ACK mechanism used by TCP) or an implicit handshake mechanism that occurs during the first exchange of application messages. Once a connection has been established it remains active until it is terminated by one of the ports, or until the communication path between the ports is severed (for example, by the failure of the node on which one of the ports is running); TIPC then immediately notifies the affected port(s) that the connection has terminated.

Connectionless messaging allows a port to exchange messages with one or more ports elsewhere in the network. A given message can be sent to a single port (unicast) or to a collection of ports (multicast), depending on the destination address specified when the message is sent.

TIPC is designed to be a reliable messaging mechanism, in which an application can send a message and assume that the message will be delivered to the specified destination as long as that destination is reachable. If a message cannot be delivered the message sender can specify whether it should be returned to its point of origin or discarded, according to the needs of the application. (See the section on "Message Delivery" that appears later in this document for more information.)

#### THINGS TO REMEMBER:

- In a TIPC network where different nodes may be running on different CPU types and/or operating systems applications must ensure that the internal structure of a message is well-defined, and accounts for any differences in message content endianness, field size, and field padding.

### 1.3 TIPC Addressing

-----

TIPC uses 3 distinct forms of addressing within a network.

#### 1.3.1 Network Address

-----

The network address was introduced in section 1.1, and is typically denoted as <Z.C.N>. Applications use this address format with certain operations to specify the portion of a TIPC network the operation applies to:

- a) <Z.C.N> indicates a network node
- b) <Z.C.0> indicates a network cluster
- c) <Z.0.0> indicates a network zone
- d) <0.0.0> has special meaning, which is operation-specific

When coding, a network address is represented as an unsigned 32-bit value, comprising 3 fields: an 8-bit zone field, a 12-bit cluster field, and a 12-bit node field. (See section 1.6 for a description of routines for constructing and deconstructing networks addresses.)

#### 1.3.2 Port Identifier

-----

Each port in a TIPC network has a unique "port identifier" or "port ID", which is typically denoted as <Z.C.N:ref>. The port ID is assigned automatically by TIPC when the port is created, and consists of the 32-bit network address of the port's node and a 32-bit reference value. The reference value is guaranteed to be unique on a per-node basis and will not be reused for a long time once the port ceases to exist.

#### 1.3.3 Port Naming

-----

While a TIPC port can send messages to another port by specifying the port ID of the destination port, it is usually more convenient to use a "functional address" that does not require the sending port to know the physical location of the destination within the network. This simplifies communication when server ports are being created, deleted, or relocated dynamically, or when multiple ports are providing a given service.

The basic unit of functional addressing within TIPC is the "port name", which is typically denoted as {type,instance}. A port name consists of a 32-bit type field and a 32-bit instance field, both of which are chosen by the application. Typically, the type field is used to indicate the class of service provided by the port, while the instance field can be used as a sub-

class indicator.

Unlike port IDs, port names need not be unique within a TIPC network. Applications are permitted to assign a given port name to multiple ports, and to assign multiple port names to a given port, or both. Port names can also be unbound from a port if they are no longer required.

Whenever an application binds a port name to a port, it must specify the level of visibility, or "scope", that the name has within the TIPC network: either "node scope", "cluster scope", or "zone scope". TIPC then ensures that only applications within that portion of the network (i.e. the same node, the same cluster, or the same zone) can access the port using that name.

To simplify the task of specifying a range of similar port name instances TIPC supports the concept of the "port name sequence", which is typically denoted as {type, lower bound, upper bound}. A port name sequence consists of a 32-bit type field and a pair of 32-bit instance fields, and represents the set of port names from {type, lower bound} through {type, upper bound}, inclusive. The lower bound of a name sequence cannot be larger than the upper bound.

There are a number of restrictions on port naming that programmers need to be aware of.

- 1) The type values from 0 to 63 are reserved by TIPC and cannot be used to designate an application service.
- 2) Port names and name sequences are designed for use by server ports. TIPC does not allow a named sever port to initiate a connection (as if it were a client port), nor does it allow the assignment of names to a connected client port (as if it were a server port).
- 3) TIPC does not currently allow the creation of partially overlapping port name sequences (unless the name sequences cannot be seen simultaneously). For example, once a node has a port having the name sequence {100, 1000, 2000} -- or the node is notified that another node has a port with this name sequence -- it cannot then assign the name sequences {100, 500, 1200} or {100, 1100, 1500} to any of its ports; however, the node is permitted to assign {100, 1000, 2000} to other ports or to use name sequences having other type values such as {150, 500, 1200}. Overlapping name sequences *are* permitted if they are published by different nodes and are published with non-overlapping scopes; for example, you can publish {100, 500, 1200} on node <1.1.1> and {100, 1100, 1500} on node <1.1.2> as long as they both are published with node scope.

#### THINGS TO REMEMBER:

- Programmers typically only have to worry about the selection of TIPC names and name sequences. (Network addresses are chosen by the TIPC network administrator, while port ID's are chosen by TIPC automatically.)
- Well-known names (or name sequences) are used in a TIPC network in the same way that well-known port numbers are used in an IP network.
- An application must use TIPC names (or name sequences) that do not conflict with the names used by other applications.

### 1.4 Using Port Names

---

This section discusses more advanced aspects of TIPC's functional addressing.

#### 1.4.1 Address Resolution

-----  
Whenever an application specifies a port name as the destination address of a message, it must also indicate where within the network TIPC should look to find the destination by specifying a "lookup domain".

The most commonly used lookup domain is <0.0.0>, which tells TIPC to use a "closest first" approach. TIPC first looks on the sending node to find a port having the specified port name; if more than one such port exists, TIPC selects one in a round-robin manner. If the sending node does not contain a matching port, TIPC then looks to all other nodes in the sending node's cluster to see if any ports have published that name using cluster scope or zone scope; again, if more than one such port exists, TIPC selects one in a round-robin manner. Finally, if no matching port is found within the sending node's cluster, TIPC looks at all other nodes in the sending node's zone for ports with a matching name and having zone scope, and selects one in a round-robin manner. (In short, address resolution is performed using 3 lookup domains in succession: first using <Z.C.N>, then <Z.C.O>, and finally <Z.O.O>.) This algorithm results in the message being delivered to a suitable destination as quickly as possible, and also load sharing similarly named messages among all such destinations at the same distance from the sender.

Alternatively, an application can specify a single lookup domain to be used for address resolution. Specifying a lookup domain of the form <Z.C.N>, <Z.C.O>, or <Z.O.O> tells TIPC to take all ports with compatible name and scope values within the specified node, cluster, or zone, respectively, and then select one in a round-robin manner. These forms can be useful in preventing a message from being sent off-node, or for evenly distributing work to all servers scattered throughout a cluster or zone.

It should be noted that the round-robin selection mechanism used by TIPC is shared by all applications using a given node. So, for example, if there are two ports within the specified lookup domain that have the desired port name, an application cannot assume that two successive messages it sends to that name will be distributed one to each port. This is because a similarly named message sent by another application may arrive in the interim, thereby causing the second message sent by the first application to go to the same destination as its first message.

#### 1.4.2 Multicast Messaging

-----

Whenever an application specifies a port name sequence as the destination address of a message (rather than a port name), this instructs TIPC to send a copy of the message to every port in the sender's cluster that has at least one port name within the destination name sequence.

This is most easily illustrated using an example. Suppose a multicast message is sent to {1000,100,200}, then the following ports will each receive exactly one copy of the message:

<1.1.10:1234> having {1000,100}	- one matching name
<1.1.11:4321> having {1000,123} and {1000,175}	- two matching names
<1.1.10:5678> having {1000,150} and {2000,150}	- non-matching name ignored
<1.1.12:5555> having {1000,110,120}	- subset overlap
<1.1.10:8888> having {1000,50,500}	- superset overlap
<1.1.14:9999> having {1000,170,300}	- partial overlap

while the following ports will not receive a copy at all:

<1.1.10:1111> having {2000, 100, 200}	- name type mismatch
<1.1.10:4444> having {1000, 50, 75}	- no overlap
<1.1.10:6666> having no names bound to it	- no overlap

Note that a port never receives more than one copy of the multicast message, even if it has several port names (or port name sequences) bound to it that overlap the specified destination name sequence.

Also note that the requirement that the destination address for a multicast message be a name sequence does not prevent applications from multicasting to a single port name; an application can simply specify a name sequence that encompasses a single instance value, such as {1000, 123, 123}.

#### THINGS TO REMEMBER:

- Multicast messaging can only be done in a connectionless manner, as TIPC does not support the concept of a "one to many" or "many to many" connection.
- It is not possible to limit the distribution of a multicast message to the ports within a given node by specifying a "lookup domain", as can be done with unicast messages.

#### 1.4.3 Name Subscriptions

TIPC provides a network topology service that applications can use to receive information about what port names exist within the application's network zone.

An application accesses the topology service by opening a message-based connection to port name {1,1} and then sending "subscription" messages to the topology service that indicate the port names of interest to the application; in return, the topology service sends "event" messages to the application when these names are published or withdrawn by ports within the network. Applications are allowed to have multiple subscriptions active at the same time; issuing a new subscription does not affect any existing subscription.

A subscription request message must contain the following information:

- 1) The port name sequence of interest to the application.

Applications that are interested in a single port name can specify a port name sequence in which the lower and upper instance values are the same.

- 2) An event filter specifying which events are of interest to the application.

The value TIPC\_SUB\_PORTS causes the topology service to generate a TIPC\_PUBLISHED event for each port name or port name sequence it finds that overlaps the specified port name sequence; a TIPC\_WITHDRAWN event is issued each time a previously reported name becomes unavailable. The value TIPC\_SUB\_SERVICE causes the topology service to generate a single publish event for the first port it finds with an overlapping name and a single withdraw event when the last such port becomes unavailable. Thus, the latter event filter allows the topology service to inform the application if there are *\*any\** ports of interest, while the former informs it about *\*all\** such ports.

- 3) A subscription timeout value.

If the subscription is still active after the specified number of milliseconds, a TIPC\_SUBSCR\_TIMEOUT event message is sent to the application and the topology service deletes the subscription. (The value

TIPC\_WAIT\_FOREVER can be specified if no time limit is desired.)

- 4) An 8 byte "user handle" that is application-defined.

This value is returned to the application as part of all events associated with the subscription request. Applications may find it useful to use this field to hold a unique subscription identifier when multiple subscription requests are active simultaneously.

An event message contains the following information:

- 1) A code indicating the type of event that has occurred.

This may be either TIPC\_PUBLISHED, TIPC\_WITHDRAWN, or TIPC\_SUBSCR\_TIMEOUT.

- 2) The instance values denoting the lower and upper bounds of the port name sequence that overlaps the name sequence specified by the subscription.

The name type value is not supplied as it is always equal to the value specified by the subscription request.

- 3) The port ID of the associated port.

- 4) The subscription request associated with the event.

The exchange of messages between application and topology service is entirely asynchronous. The application may issue new subscription requests at any time, while the topology service may send event messages about these subscriptions to the application at any time.

The connection between the application and the topology service continues until the application terminates it, or until the topology service encounters an error that requires it to terminate the connection. When the connection ends, any active subscription requests are automatically cancelled by TIPC.

#### THINGS TO REMEMBER:

- It is not possible to limit the range of a subscription request to a specific node, cluster, or zone by specifying a lookup domain; the topology service always monitors the requestor's entire zone for matching port names.
- Every node in a TIPC network automatically publishes a port name of the form {0,<Z.C.N>}, where <Z.C.N> is the node's network address. Applications can determine what nodes currently comprise the network, and track the subsequent arrival and departure of nodes from the network, by creating a subscription that tracks the publication and withdrawal of names for type 0. The dummy subscription example in section 5.1 below illustrates this technique.

#### CURRENT LIMITATIONS:

- The only way to cancel a subscription (other than letting it time out) is to close the connection to the topology service, thereby cancelling all subscriptions issued on that connection.

### 1.5 Message Delivery

-----

On the surface, message delivery in TIPC is a simple series of steps: a message is created by a sender, TIPC carries it to the specified destination, and the receiver consumes the message. And, in practice, this is exactly what happens most of the time. However, there are a number of places along the way where things can get complicated, and in these cases it is important for application

designers to understand exactly what TIPC will do.

The sections that follow describe the various steps performed by TIPC during the exchange of a unicast message; the final section outlines how any differences that occur when dealing with a multicast message.

#### 1.5.1 Message Creation

---

The first step in sending a message is to create it. The most common reason TIPC is unable to create a message is because the sender passes in one or more invalid arguments to the send routine. The term "invalid" refers both to values that are never acceptable under any circumstances (such as specifying a message length greater than 66000 bytes) and to values that are not acceptable for the current sender (such as requesting a send operation on a socket that has been turned into a listening socket).

Other reasons that TIPC may be unable to create a message:

- There are no more message buffers available that TIPC can use.
- The link TIPC selected to carry the message to its destination was congested and the sender did not want to block until the congestion cleared (see 1.5.3 below).
- The peer socket on a connection was congested (i.e. had too many unconsumed messages in its receive queue) and the sender did not want to block until the congestion cleared (see 1.5.6 below).

In all of these cases the send operation will return a failure code indicating that the intended message was not sent. If the message is created successfully the send operation returns a success indication.

#### THINGS TO REMEMBER:

- If the sender specifies a destination address that does not currently exist within the TIPC network, TIPC does *\*NOT\** treat this as an invalid send request (i.e. it's not the sender's fault that the destination doesn't exist). Instead TIPC creates the message and then "rejects" it because it is undeliverable (see 1.5.6 below). The return value for the send operation will indicate success since the message was successfully created and processed by TIPC.

#### 1.5.2 Source Routing

---

Once a message has been created, TIPC then determines what node the message should be sent to. If the specified destination address is a port ID, the destination node is pre-determined; if the address is a port name, TIPC performs a name table lookup to select a port (see 1.4.1 above), and then uses the node associated with that port. The message is then passed to a link for off-node transmission (see 1.5.3 below) or is handed off to the destination port directly if it is on the same node as the sender (see 1.5.5 below).

Problems that can arise during the source routine phase of message delivery:

- No matching port can be located during a name table lookup when sending by port name.
- No working link to the specified destination node can be found when sending by port ID.

In all cases of source routing failure, the message is rejected (see section



1.5.6 below).

#### THINGS TO REMEMBER:

- A message that specifies a port name and is sent off-node may not actually end up going to the port selected during the name table lookup, since the destination node will perform a second name table lookup when it receives the message (see 1.5.4 below).

#### 1.5.3 Link Transmission

-----

Once a message is given to a link for transmission to another node, the link will normally deliver the message to that node even if problems arise. For example, the link will automatically detect lost messages and retransmit them, or will re-route messages over an alternate link if it loses contact with its peer link endpoint on the other node. Such error recovery is possible because TIPC keeps a copy of each outgoing message in a transmit queue until it is notified that the message has been successfully received by the peer link endpoint.

If a link endpoint's transmit queue grows too large because the peer link endpoint falls behind in acknowledging the successful arrival of messages (typically around 50 messages), TIPC declares "link congestion" on that link. When a link becomes congested, the link only accepts a new message for transmission if it is important enough (i.e. the more important the message, the longer the queue is allowed to be).

Whenever a message cannot be sent because of link congestion, TIPC checks the "source droppable" setting of the sending port. If the setting is enabled (indicating that the message is being sent in an unreliable manner) TIPC discards the message, but provides no indication of this to the sender. If the source droppable setting is disabled (which is the default case), TIPC will normally block the sending application until the congestion clears, and then resume the send operation; however, if the application has requested a non-blocking send, the application will not block when link congestion occurs and the send operation returns a failure indication.

In the event that a link to a destination node fails and there are no other links available that can be used to re-route traffic, any messages in the link's transmit queue are simply discarded. The messages are *\*not\** rejected (and potentially returned to their originating ports) because TIPC does not know whether or not they were successfully delivered.

#### 1.5.4 Destination Routing

-----

Once a message arrives at the specified destination node over a link, TIPC then determines what port it should be sent to on that node.

If the specified destination address is a port ID, the destination port is pre-determined; if no such port exists the message is considered undeliverable and rejected (see 1.5.6 below).

If the destination address is a port name, TIPC performs a name table lookup and selects a port (see 1.4.1 above). If no such port exists TIPC repeats the source routing operation and tries to send the message to another node; if no such node can be found, or if the message has been previously re-routed too many times, the message is considered undeliverable and rejected (see 1.5.6 below).

### 1.5.5 Message Consumption

---

When a message (finally!) reaches the destination port it is either consumed immediately (if the controlling application is using the native API) or added to a receive queue (if the controlling application is using the socket API). In the latter case, the message typically remains in the socket's receive queue until it is received by the application that owns the socket. Queued messages are consumed by the application in a FIFO manner, and once the contents of a message have been passed to the application the message is discarded.

If an application terminates access to the socket (using either the `close()` or `shutdown()` APIs) before all messages in the receive queue are consumed, all unconsumed messages are considered undeliverable and are rejected (see 1.5.6 below).

It is very important that TIPC applications be engineered to consume their incoming messages at a rate that prevents them from accumulating in large numbers in any socket receive queue. Failure to do so can result in TIPC declaring either "port congestion" or "socket congestion".

Port congestion can occur once more than 512 messages have accumulated in the receive queue of a connection-oriented socket. Once this is detected, TIPC may block the peer socket from sending messages until the congestion clears (see 1.5.1 above) or, if the sender is sending in an unreliable manner, cause such messages to be discarded (see 1.5.3 above).

Socket congestion can occur once TIPC detects that too many unreceived messages exist on a node or on an individual socket. More precisely, a node can have up to 5000 messages sitting in socket receive queues before congestion handling kicks in; once this happens, low importance messages will be rejected (see 1.5.6 below) but higher importance messages will continue to be accepted. Medium importance messages get screened out once the number of pending messages hits 10000, and high priority messages at 500,000 messages; critical priority messages are always accepted. Similar congestion handling occurs on a per-socket basis, but the thresholds are one half the global threshold values (i.e. at 2500, 5000, and 250,000).

Since the impact of socket congestion is more significant for a connection-oriented socket than port congestion (i.e. it terminates the connection), the smaller port congestion threshold has been chosen so that it will normally kick in first and prevent the socket receive queue from growing larger. However, the existence of the per-node socket congestion threshold means that it is possible for socket congestion to occur before port congestion occurs.

NOTE: These message congestion thresholds may be more configurable in future releases of TIPC since it's not really realistic to have a one-size-fits-all solution that will work well on a wide variety of hardware configurations (i.e. a resource-constrained DSP will probably need lower thresholds than a resource-rich Linux box).

### 1.5.6 Message Rejection

---

When a message is "rejected" because it cannot be delivered, TIPC checks the message's "destination droppable" setting to see what the sender wanted done with the message.

If the destination droppable setting is enabled, TIPC simply discards the message. This setting is the default for messages sent using a connectionless

socket, and was chosen to simplify the job of porting applications written for UDP to use TIPC.

If the destination droppable setting is disabled, TIPC sends the first 1024 bytes of the message back to the message originator; such a message is called a "returned message". An error code is also incorporated into each returned message to allow the sender to determine why the message was returned. In the case of a connection-oriented message, the return of an undeliverable message also causes the connection to be terminated at both ends. By default, the destination droppable setting is disabled for messages sent using connection-oriented sockets; this decision was made to simplify the job of porting applications written for TCP to use TIPC.

TIPC's socket API has been designed so that applications that don't want to concern themselves with returned messages can easily ignore them. However, the ability for a sending application to examine returned messages can be helpful in debugging problems during the design and testing of a new TIPC application.

#### THINGS TO REMEMBER:

- The returned message capability of TIPC must NOT be used by a sending application to determine what messages were successfully consumed by the receiving application! While the return of a message indicates that the receiver did not consume the message, the non-return of a message does not indicate that it was successfully consumed. (For example, if a destination node suffers a power failure, TIPC will be unable to return any messages that are sitting unprocessed in a socket receive queue.) The only way for the sending application to know that a message was consumed is for it to receive an explicit acknowledgement message generated by the receiving application.

#### 1.5.7 Multicast Message Delivery

---

Multicast message creation is done the same way as for unicast messages, but since multicasting is always done in a connectionless manner it is not possible for peer port congestion to occur.

Multicast source routing always involves a name table lookup of a port name sequence. If no port within the cluster overlaps the specified name sequence the message is simply discarded. Otherwise, TIPC sends a copy of the message to each overlapping port on the sending node and also determines if any off-node ports have an overlap; if there is at least one such port then the message is passed to a special multicast link.

The multicast link operates much like a regular unicast link, except that it sends its messages to *\*all\** nodes in the cluster rather than just one, and has a smaller congestion threshold (around 20 messages).

Whenever the multicast link delivers the message to a node, TIPC repeats the name table lookup and sends a copy of the message to all overlapping ports it finds on that node; if there are no such ports the message is discarded. Once a multicast message arrives at a destination port, it is treated just like a unicast message and is subject to the same socket congestion and message rejection handling.

#### THINGS TO REMEMBER:

- TIPC currently does not permit an application to send a multicast message with the "destination droppable" setting disabled; consequently, TIPC will never try to return an undeliverable multicast message to its sender.

## 1.6 Routines

The following utility routines are available to programmers:

`tipc_addr()` - combine zone, cluster, and node numbers into a TIPC address  
`tipc_cluster()` - take a TIPC network address and return the cluster number  
`tipc_node()` - take a TIPC address and return the node number  
`tipc_zone()` - take a TIPC address and return the zone number

Further information about each of these routines is provided in the following sections.

### 1.6.1 `tipc_addr`

`u32 tipc_addr(unsigned int zone, unsigned int cluster, unsigned int node)`

This routine takes individual zone, cluster, and node numbers and combines them into a 32-bit TIPC network address.

### 1.6.2 `tipc_cluster`

`unsigned int tipc_cluster(u32 addr)`

This routine takes a 32-bit TIPC network address and returns the cluster number contained in the address.

### 1.6.3 `tipc_node`

`unsigned int tipc_node(u32 addr)`

This routine takes a 32-bit TIPC network address and returns the node number contained in the address.

### 1.6.4 `tipc_zone`

`unsigned int tipc_zone(u32 addr)`

This routine takes a 32-bit TIPC network address and returns the zone number contained in the address.

## 2. Socket API

The TIPC socket API allows programmers to access the capabilities of TIPC using the well-known socket paradigm.

### IMPORTANT:

TIPC does not support all socket API routines available with other socket-based protocols, nor does it support all possible capabilities for the routines that are provided. Likewise, certain new capabilities provided by TIPC have been made available by adapting the socket API to accommodate them.

Programmers who have used the socket API with other protocols are strongly advised to read this section carefully to ensure they understand where the TIPC socket API differs from their previous experiences.

### GENERAL LIMITATIONS:

TIPC's socket API is currently tailored for use by single-threaded applications; consequently, if multiple threads of control try to perform I/O operations on a given socket it is possible that some threads may become blocked unexpectedly. Most significantly, if there is a thread blocked trying to read from a socket, any other thread that writes to the socket will block until the read operation is completed. Until this situation is rectified, multi-threaded applications should use `select()` or `poll()` to test whether a socket is ready for reading before beginning a blocking read operation.

## 2.1 Routines

The following socket API routines are available to programmers:

<code>accept( )</code>	- accept a new connection on a socket
<code>bind( )</code>	- bind or unbind a TIPC name to the socket
<code>close( )</code>	- close the socket
<code>connect( )</code>	- connect the socket
<code>getpeername( )</code>	- get the port ID of the peer socket
<code>getsockname( )</code>	- get the port ID of the socket
<code>getsockopt( )</code>	- get the value of an option for the socket
<code>listen( )</code>	- listen for socket connections
<code>poll( )</code>	- input/output multiplexing
<code>recv( )</code>	- receive a message from the socket
<code>recvfrom( )</code>	- receive a message from the socket
<code>recvmsg( )</code>	- receive a message from the socket
<code>send( )</code>	- send a message on the socket
<code>sendmsg( )</code>	- send a message on the socket
<code>sendto( )</code>	- send a message on the socket
<code>setsockopt( )</code>	- set the value of an option for the socket
<code>shutdown( )</code>	- shut down socket send and receive operations
<code>socket( )</code>	- create an endpoint for communication

Further information about each of these routines can be found in the following sections.

### 2.1.1 accept

```
int accept(int sockfd, struct sockaddr *cliaddr, socklen_t *addrlen)
```

Accepts a new connection on a socket.

Comments:

- If non-NULL, "cliaddr" is set to the port ID of the peer socket. This info can be used to perform basic validation of the connection requestor's identity (eg. disallow connections originating from certain network nodes).

### 2.1.2 bind

```
int bind(int sockfd, const struct sockaddr *myaddr, socklen_t addrlen)
```

Binds or unbinds a TIPC name (or name sequence) to the socket. A bind operation is requested by setting "myaddr->scope" to `TIPC_NODE_SCOPE`, `TIPC_CLUSTER_SCOPE`, or `TIPC_ZONE_SCOPE`, as appropriate. An unbind operation is requested by setting "myaddr->scope" to arithmetic inverse of the scope used when the name was bound (eg. `-TIPC_NODE_SCOPE`). Specifying zero for "addrlen" unbinds all names and name sequences currently bound to the socket.

Comments:

- It is legal to bind more than one TIPC name or name sequence to a socket.
- If a socket is currently connected to a peer it is considered to be unavailable to serve other clients and cannot use bind() to bind an additional TIPC name to itself.
- bind() cannot be used to changed the port ID of a socket.

### 2.1.3 close

```
int close(int)
```

Closes the socket.

Comments:

- Any unprocessed messages remaining in the socket's receive queue are rejected (i.e. discarded or returned), as appropriate.

### 2.1.4 connect

```
int connect(int sockfd, const struct sockaddr *servaddr, socklen_t addrlen)
```

Attempts to make a connection on the socket using TIPC's explicit handshake mechanism.

Comments:

- If "servaddr" is set to a TIPC name (but not a TIPC name sequence or a TIPC port ID), the "addr.name.domain" field can be used to affect the name lookup process. The "scope" field of "servaddr" is ignored by connect().
- If a socket has a name bound to it, it is considered to be a server and cannot use connect() to initiate a connection as a client.
- TIPC does not support the use of connect() with connectionless sockets. (POSIX non-conformity)
- TIPC does not support the non-blocking form of connect(); use sendto() to establish connections using implicit handshaking to achieve this effect. (POSIX non-conformity)

### 2.1.5 getpeername

```
int getpeername(int sockfd, struct sockaddr *peeraddr, socklen_t *addrlen)
```

Gets the port ID of the peer socket.

Comments:

- The use of "name" in getpeername() can be confusing, as the routine does not actually return the TIPC names or name sequences that have been bound to the peer socket.

### 2.1.6 getsockname

```
int getsockname(int sockfd, struct sockaddr *localaddr, socklen_t *addrlen)
```

Gets the port ID of the socket.

Comments:

- The use of "name" in getsockname() can be confusing, as the routine does not actually return the TIPC names or name sequences that have been bound to the socket.

### 2.1.7 getsockopt

-----  
int getsockopt(int sockfd, int level, int optname,  
void \*optval, socklen\_t \*optlen)

Gets the current value of a socket option. For a description of the options supported when "level" is SOL\_TIPC, see the description for setsockopt() below.

Comments:

- TIPC does not currently support socket options for level SOL\_SOCKET, such as SO\_SNDBUF.
- TIPC does not currently support socket options for level IPPROTO\_TCP, such as TCP\_MAXSEG. Attempting to get the value of these options on a SOCK\_STREAM socket returns the value 0.

#### 2.1.8 listen

-----

int listen(int sockfd, int backlog)

Enables a socket to listen for connection requests.

Comments:

- The "backlog" parameter is currently ignored.

#### 2.1.9 poll

-----

int poll(struct pollfd \*fdarray, unsigned long nfds, int timeout)

Indicates the readiness of the specified TIPC sockets for I/O operations, using the standard poll() mechanism.

TIPC currently sets the returned event flags as follows:

- POLLRDNORM and POLLIN are set when the socket's receive queue is non-empty. (They are also set for a connection-oriented, non-listening socket whose connection has been terminated or has not yet been initiated, since a receive operation on such a socket will fail immediately.)
- POLLOUT is set except when a socket's connection has been terminated.
- POLLHUP is set when a socket's connection has been terminated.

Comments:

- It is important to realize that the poll() bits indicate that the associated input/output operation will not block, NOT that the operation will be successful!
- The POLLOUT event flag cannot be used in isolation to guarantee that a send operation performed on the socket will not block. Since outgoing messages are queued on a per-link basis rather than a per-socket basis, sending to a destination that is routed through link A may be blocked while sending to a destination that is routed through link B may not be blocked. To ensure that an application does not blocked during a send operation, it should use the MSG\_DONTWAIT flag or set the socket for nonblocking I/O using fcntl().
- A socket that is connecting asynchronously is considered writeable, since attempting a second send operation during an implied connection setup will immediately fail. (POSIX non-conformity)

#### 2.1.10 recv

-----

ssize\_t recv(int sockfd, void \*buff, size\_t nbytes, int flags)

Attempts to receive a message from the socket.

Comments:

- When used with a connectionless socket, a return value of 0 indicates the return of an undelivered data message that was originally sent by this socket.
- When used with a connection-oriented socket, a return value of 0 or -1 indicates connection termination. A value of 0 indicates that the connection was terminated by the peer using shutdown(); connection termination by any other means causes a return value of -1.
- Applications can determine the exact cause of connection termination and/or message non-delivery by using recvmsg() instead of recv().
- TIPC supports the MSG\_PEEK flag when receiving, as well as the MSG\_WAITALL flag when receiving on a SOCK\_STREAM socket; all other flags are ignored.

#### 2.1.11 recvfrom

```
ssize_t recvfrom(int sockfd, void *buff, size_t nbytes, int flags,  
                 struct sockaddr *from, socklen_t *addrlen)
```

Attempts to receive a message from the socket. If successful, the port ID of the message sender is returned in "from".

Comments:

- See the comments section for recv().

#### 2.1.12 recvmsg

```
ssize_t recvmsg(int sockfd, struct msghdr *msg, int flags)
```

Attempts to receive a message from the socket. If successful, the port ID of the message sender is captured in the "msg\_name" field of "msg" (if non-NULL) and ancillary data relating to the message is captured in the "msg\_control" field of "msg" (if non-NULL).

The following ancillary data objects may be captured:

- 1) TIPC\_ERRINFO - The TIPC error code associated with a returned data message or a connection termination message, and the length of the returned data. (8 bytes: error code + data length)
- 2) TIPC\_RETDATA - The contents of a returned data message, up to a maximum of 1024 bytes.
- 3) TIPC\_DESTNAME - The TIPC name or name sequence that was specified by the sender of the message. (12 bytes: type + lower instance + upper instance; the latter two values are the same for a TIPC name, but may differ for a name sequence)

Each of these objects is only created where relevant. For example, receipt of a normal data message never creates the TIPC\_ERRINFO and TIPC\_RETDATA objects, and only creates the TIPC\_DESTNAME object if the message was sent using a TIPC name or name sequence as the destination rather than a TIPC port ID. Those objects that are created will always appear in the relative order shown above.

If ancillary data objects capture is requested (i.e. "msg->msg\_control" is non-NULL) but insufficient space is provided, the MSG\_CTRUNC flag is set to indicate that one or more available objects were not captured.



Comments:

- When used with a connectionless socket, a return value of 0 indicates the arrival of a returned data message that was originally sent by this socket.
- When used with a connection-oriented socket, a return value of 0 or -1 indicates connection termination. The exact return value upon connection termination is influenced by the "msg\_control" field of "msg". If "msg\_control" is NULL, a return value of 0 indicates that the connection was terminated by the peer using shutdown(); connection termination by any other means causes a return value of -1. If "msg\_control" is non-NULL, a return value of 0 is always used; the application must examine the TIPC\_ERRINFO object to determine if the connection was explicitly terminated by the peer. (POSIX non-conformity)
- When used with connection-oriented sockets, TIPC\_DESTNAME is captured for each data message received by the socket if the connection was established using a TIPC name or name sequence as the destination address. Note: There is currently no way for the destination socket to capture TIPC\_DESTNAME following accept() until the originator sends a data message.
- TIPC supports the MSG\_PEEK flag when receiving, as well as the MSG\_WAITALL flag when receiving on a SOCK\_STREAM socket; all other flags are ignored.

#### 2.1.13 send

-----

```
ssize_t send(int sockfd, const void *buff, size_t nbytes, int flags)
```

Attempts to send a message from the socket to its peer socket.

Comments:

- send() should not be used until a connection has been fully established using either explicit or implicit handshaking.
- TIPC supports the MSG\_DONTWAIT flag when sending; all other flags are ignored.

#### 2.1.14 sendmsg

-----

```
ssize_t sendmsg(int sockfd, struct msghdr *msg, int flags)
```

Attempts to send a message from the socket to the specified destination. If the destination is denoted by a TIPC name or a port ID the message is unicast to a single port; if the destination is denoted by a TIPC name sequence the message is multicast to all ports having a TIPC name or name sequence that overlaps the destination name sequence.

Comments:

- See the comments section for sendto().
- TIPC does not currently support the use of ancillary data with sendmsg().

#### 2.1.15 sendto

-----

```
ssize_t sendto(int sockfd, const void *buff, size_t nbytes, int flags,  
               const struct sockaddr *to, socklen_t addrlen)
```

Attempts to send a message from the socket to the specified destination. If the destination is denoted by a TIPC name or a port ID the message is unicast to a single port; if the destination is denoted by a TIPC name sequence the message is multicast to all ports having a TIPC name or name sequence that overlaps the destination name sequence.

Comments:

- If the destination address is a TIPC name the "addr.name.domain" field indicates the search domain used during the name lookup process. (In contrast, if the destination address is a TIPC name sequence the default "closest first" algorithm is always used; if it is a TIPC port ID no name lookup occurs.) The "scope" field of the destination address is always ignored when sending.
- TIPC supports the MSG\_DONTWAIT flag when sending; all other flags are ignored.
- A connection-oriented socket that is unconnected can initiate connection establishment using implicit handshaking by simply sending a message to a specified destination, rather than using connect(). However, the connection is not fully established until the socket successfully receives a message sent by the destination using recv(), recvfrom(), or recvmsg().

#### 2.1.16 setsockopt

```
int setsockopt(int sockfd, int level, int optname,  
               const void *optval, socklen_t optlen)
```

Sets a socket option to the specified value. Currently, the following values of "optname" are supported when "level" is SOL\_TIPC:

##### 1) TIPC\_IMPORTANCE

This option governs how likely a message sent by the socket is to be affected by congestion. A message with higher importance is less likely to be delayed or dropped due to link congestion, and also less likely to be rejected due to receiver congestion. The following values are defined: TIPC\_LOW\_IMPORTANCE, TIPC\_MEDIUM\_IMPORTANCE, TIPC\_HIGH\_IMPORTANCE, and TIPC\_CRITICAL\_IMPORTANCE.

By default, TIPC\_LOW\_IMPORTANCE is used for all TIPC socket types.

##### 2) TIPC\_SRC\_DROPPABLE

This option governs the handling of messages sent by the socket if link congestion occurs. If enabled, the message is discarded; otherwise the system queues the message for later transmission.

By default, this option is disabled for SOCK\_SEQPACKET, SOCK\_STREAM, and SOCK\_RDM socket types (resulting in "reliable" data transfer), and enabled for SOCK\_DGRAM (resulting in "unreliable" data transfer).

##### 3) TIPC\_DEST\_DROPPABLE

This option governs the handling of messages sent by the socket if the message cannot be delivered to its destination, either because the receiver is congested or because the specified receiver does not exist. If enabled, the message is discarded; otherwise the message is returned to the sender.

By default, this option is disabled for SOCK\_SEQPACKET and SOCK\_STREAM socket types, and enabled for SOCK\_RDM and SOCK\_DGRAM. This arrangement ensures proper teardown of failed connections when connection-oriented data transfer is used, without increasing the complexity of connectionless data transfer.

##### 4) TIPC\_CONN\_TIMEOUT

This option specifies the number of milliseconds connect() will wait before aborting a connection attempt because the destination has not responded. By default, 8000 (i.e. 8 seconds) is used.

This option has no effect when establishing connections using sendto().

Comments:

- TIPC does not currently support socket options for level SOL\_SOCKET, such as SO\_SNDBUF.
- TIPC does not currently support socket options for level IPPROTO\_TCP, such as TCP\_MAXSEG. Setting these options on a SOCK\_STREAM socket has no effect.

#### 2.1.17 shutdown

-----  
int shutdown(int sockfd, int howto)

Shuts down socket send and receive operations on a connection-oriented socket. The socket's peer is notified that the connection was deliberately terminated by the application (by means of the TIPC\_CONN\_SHUTDOWN error code), rather than as the result of an error.

Comments:

- Applications should normally call shutdown() to terminate a connection before calling close().
- TIPC does not support partial shutdown of a connection; attempting to shut down either send or receive operations always shuts down both.
- A socket that has been shutdown() cannot be re-used for a new connection; this prevents any "stale" incoming messages from an earlier connection from interfering with the new connection.

#### 2.1.18 socket

-----  
int socket(int family, int type, int protocol)

Creates an endpoint for communication.

TIPC currently supports the following values for "type":

- a) SOCK\_DGRAM - for unreliable connectionless messages
- b) SOCK\_RDM - for reliable connectionless messages
- c) SOCK\_SEQPACKET - for reliable connection-oriented messages
- d) SOCK\_STREAM - for reliable connection-oriented byte streams

Comments:

- The "family" parameter should always be set to AF\_TIPC.
- The "protocol" parameter should always be set to 0.

### 2.2 Examples

-----  
A variety of demo programs can be found at <http://sourceforge.net/projects/tipc>, which may be useful in understanding how to write an application that uses TIPC.

### 3. Native API

-----  
The TIPC native API allows programmers to access the capabilities of TIPC in a more direct manner than with the socket API.

Benefits of native API:

1. Low-level operation can lead to faster execution speed.
2. Can exclude socket code from system to reduce object code size.

Limitations of native API:

1. Not available to user-space applications.
2. Low-level operation places a greater burden on programmer.

### 3.1 Concepts

-----

There are a number of important conceptual differences between programming with the native API and programming with the socket API. Understanding these concepts is an essential pre-requisite for using the native API effectively.

#### Ports:

The fundamental communication endpoint of the native API is a "port", which operates at a much more primitive level than a socket. Applications using TIPC ports are sometimes required to deal with aspects of the TIPC protocol that were hidden by the socket API, including handling undeliverable messages that are returned to the sending port and managing the handshaking required to set up and tear down port-to-port connections.

#### Port reference:

Every TIPC port has a unique "reference" value, which is analogous to the file descriptor value that is associated with a socket. Native API routines that manipulate ports use the port reference argument to identify the port, rather than a pointer to the actual port data structure; this allows TIPC to gracefully handle cases where an application inadvertently attempts to utilize a port that no longer exists.

#### User registration:

TIPC allows an application using the native API to register as a "user", and assigns it a user identifier. If this user identifier is provided by the application when it creates a port, TIPC will delete the port automatically if the application later deregisters itself. This feature can simplify things for a programmer whose application uses a constantly changing set of ports, since TIPC takes care of deleting all ports currently in use by the application when the application terminates. (Applications not wishing to take advantage of this capability can skip the optional registration process entirely and simply create their ports anonymously using a user identifier value of 0.)

#### Sending messages:

Applications can send messages using the native API in much the same way as with the socket API. The message can be specified either as a set of one or more byte arrays (using the "iovec" structure) or as a socket buffer (using the "sk\_buff" structure), as long as it does not exceed TIPC's 66000 byte limit on message size. The latter form can improve performance by eliminating the need for TIPC to copy the data into a socket buffer, but for best results the application that creates the buffer should reserve 80 bytes of headroom to allow a TIPC message header and data link header to be prepended easily.

#### Receiving messages:

The native API does not provide any synchronous mechanism for receiving messages sent to a port. (That is, there is no equivalent of the `recv()`, `recvfrom()`, or `recvmsg()` routines that the socket API provides.) Instead, an application specifies a set of message handling callback routines when it creates a port; TIPC then invokes the appropriate routine each time a message is received by the port.

Individual callback routines may be specified to handle:

- 1) a direct message (i.e. one sent to a port ID)

- 2) a named message (i.e. one sent to a port name or name sequence)
- 3) a connection message (i.e. one sent on an established connection)
- 4) an errored direct message (i.e. a direct message that was returned)
- 5) an errored named message (i.e. a named message that was returned)
- 6) an errored connection message (i.e. a connection message that was returned)

An application only needs to supply callback routines for the messages that the port actually needs to handle. If TIPC receives a message for which no callback routine has been specified, it automatically rejects the message (or, in the case of an errored message, discards it).

Since the callback routine executes in a TIPC kernel thread, rather than one of the application's threads, the programmer must be prepared to handle any critical section issues that arise between the various threads. Alternatively, the callback routine can transfer responsibility to an application thread (as outlined in section 3.3.3 below), thereby allowing the application to emulate a synchronous receive capability of its own.

### 3.2 Routines

The native API routines listed below are available to programmers. More detail about the arguments and return value for each of these routines can be found by looking at the function prototypes in `tipc.h`. In many cases the use of the routine will be obvious. Unfortunately, a comprehensive description of each routine is not currently available. (Feel free to write one!) You can also consult the examples section below and/or the source code for each routine to learn more about what these routines do and how to use them.

WARNING! The native API is still under development at this time  
 WARNING! and has not been finalized. Expect changes in future  
 WARNING! versions of TIPC.

/\* TIPC operating mode routines \*/

`tipc_get_addr()` - get <Z.C.N> of own node  
`tipc_get_mode()` - get TIPC operating mode  
`tipc_attach()` - register application as a TIPC user  
`tipc_detach()` - deregister TIPC user & free all associated ports

/\* TIPC port manipulation routines \*/

`tipc_createport()` - create a TIPC port & generate reference  
`tipc_deleteport()` - delete a TIPC port & obsolete reference  
`tipc_ref_valid()` - determine if port reference is valid  
`tipc_ownidentity()` - get port ID of port  
`tipc_set_portimportance()` - set port traffic importance level  
`tipc_portimportance()` - get port traffic importance level  
`tipc_set_portunreliable()` - set port traffic "source droppable" setting  
`tipc_portunreliable()` - get port traffic "source droppable" setting  
`tipc_set_portunreturnable()` - set port traffic "destination droppable" setting  
`tipc_portunreturnable()` - get port traffic "destination droppable" setting  
`tipc_publish()` - bind name/name sequence to port  
`tipc_withdraw()` - unbind name/name sequence from port  
`tipc_connect2port()` - associate port with peer  
`tipc_disconnect()` - disassociate port with peer  
`tipc_shutdown()` - shut down connection to peer & disassociate  
`tipc_isconnected()` - determine if port is currently connected  
`tipc_peer()` - get port ID of peer port

```
/* TIPC messaging routines */
```

```
tipc_send()           - send iovec(s) on connection
tipc_send_buf()       - send sk_buff on connection
tipc_send2name()      - send iovec(s) to port name
tipc_send_buf2name()  - send sk_buff to port name
tipc_send2port()      - send iovec(s) to port ID
tipc_send_buf2port()  - send sk_buff to port ID
tipc_multicast()      - multicast iovec(s) to port name sequence

tipc_forward2name()   - [may be obsoleted]
tipc_forward_buf2name() - [may be obsoleted]
tipc_forward2port()   - [may be obsoleted]
tipc_forward_buf2port() - [may be obsoleted]
```

```
/* TIPC subscription routines */
```

```
tipc_ispublished()    - determines if a specific name has been published
tipc_available_nodes() - [likely to be obsoleted]
```

### 3.3 Examples

#### 3.3.1 Basic port operations

Create a port:

```
static u32 port_ref;

tipc_createport(0, NULL, TIPC_LOW_IMPORTANCE,
               NULL, NULL, NULL,
               NULL, named_msg_event, NULL,
               NULL, &port_ref);
```

Bind the name {100,123} with "cluster" scope to the port:

```
struct tipc_name_seq seq;

seq.type = 100 ;
seq.lower = 123 ;
seq.upper = 123 ;
tipc_publish(port_ref, TIPC_CLUSTER_SCOPE, &seq);
```

Process messages sent to port {100,123}:

```
/* Note: This callback routine was specified during port creation above */
```

```
static void named_msg_event(void *usr_handle,
                           u32 port_ref,
                           struct sk_buff **buf,
                           unsigned char const *data,
                           unsigned int size,
                           unsigned int importance,
                           struct tipc_portid const *orig,
                           struct tipc_name_seq const *dest)
{
    /* 'data' points to message content, 'size' indicates how much */
```

```

    printk("%s", data);

    /* can send reply message(s) back to originator, if desired */

    struct iovec my_iov;
    char reply_info[30];

    strcpy(reply_info, "here is the reply");
    my_iov.iov_base = reply_info;
    my_iov.iov_len = strlen(reply_info) + 1;
    tipc_send2port(port_ref, orig, 1, &my_iov);

    /* TIPC discards the received message upon exit */
}

```

Delete the port:

```
tipc_deleteport(port_ref);
```

### 3.3.2 TIPC user registration

---

Register TIPC user:

```

static u32 user_ref;

tipc_attach(&user_ref, NULL, NULL);

```

Create port and associate with registered TIPC user:

```

static u32 port_ref;

tipc_createport(user_ref, NULL, TIPC_LOW_IMPORTANCE,
                NULL, NULL, NULL,
                NULL, named_msg_event, NULL,
                NULL, &port_ref);

```

Deregister TIPC user (and all associated ports):

```
tipc_detach(user_ref);
```

### 3.3.3 Synchronous message receive

---

Application thread:

```

/* Initialize data structures */

struct sk_buff_head message_q;
wait_queue_head_t wait_q;

skb_queue_head_init(&message_q);
init_waitqueue_head(&wait_q);

/* Wait for messages; process & discard each one in turn */

while (1) {
    struct sk_buff *skb;

```

```

        if (wait_event_interruptible(&wait_q, (!skb_queue_empty(&message_q))))
            continue;

        skb = skb_dequeue(&message_q);

        < ... Process message as required ... >

        kfree_skb(skb);
    }

```

Callback routine converts asynchronous receive into synchronous receive:

```

static void named_msg_event(void *usr_handle,
                           u32 port_ref,
                           struct sk_buff **buf,
                           unsigned char const *data,
                           unsigned int size,
                           unsigned int importance,
                           struct tipc_portid const *orig,
                           struct tipc_name_seq const *dest)
{
    /* Add message to queue of unprocessed messages */

    skb_queue_tail(&message_q, *buf);

    /* Tell TIPC *not* to discard the received message upon exit */

    *buf = NULL;

    /* Wake up application */

    wake_up_interruptible(&wait_q);
}

```

### 3.3.4 More examples

-----  
 Demo programs utilizing the native API can be found at <http://tipc.sf.net>.

In addition, the TIPC source code itself contains a couple of sections that utilize the native API just like an application might:

#### 1) tipc\_cfg\_init() in net/tipc/config.c

This file contains the TIPC configuration service (using port name {0,<Z.C.N>}, which handles messages sent by the tipc-config application. It utilizes a very simple connectionless request-and-reply approach to messaging.

#### 2) tipc\_subscr\_start() in net/tipc/subscr.c

This file contains the TIPC topology service (using port name {1,1}), which handles subscription requests from applications and returns subscription events. It demonstrates the correct way to handle connection establishment (both explicit and implied) and tear down (both self-initiated and peer-initiated).

## 4. FAQ

-----  
 This section contains the answers to frequently asked questions.



#### 4.1 How can I determine what node a socket is running on?

---

Use `getsockname()` to determine the port ID of the socket, then examine the "node" field to determine the network address its node. For example, if `getsockname` returns a port ID of `<1.1.19:1234567>` then the node field will have the value `0x01001013` (representing network address `<1.1.19>`).

#### 4.2 How can I cancel a subscription?

---

The only way to cancel a subscription (other than letting it time out) is to close the connection to the topology service, thereby cancelling all subscriptions issued on that connection. The ability to cancel a single subscription will be added in an upcoming release.

#### 4.3 When should I use an implied connect instead of an explicit connect?

---

The simplest approach is to assume that you will be using an explicit connect and design your code accordingly. If you end up with a `connect()` followed by a send routine followed by a receive routine, then you should be able to combine the `connect()` and send routine into a `sendto()`, thereby saving yourself the overhead of an additional system call and the exchange of empty handshaking messages during connection establishment. On the other hand, if you end up with a `connect()` followed by a receive, or a `connect()` followed by two sends, then you can't use the implied connect approach.

### 5. Tips and Techniques

---

This section illustrates some techniques for using TIPC that may be of interest to programmers when designing applications using TIPC.

#### 5.1 Dummy Subscriptions

---

It is sometimes useful to issue a "dummy" subscription to the TIPC topology server -- that is, a subscription that has a time limit but will never match any published name. Such a subscription will never generate any publish or withdraw events, but will generate a timeout event when it expires.

For example, suppose an application wishes to report what nodes are present in the network every minute. This can be accomplished as follows:

```
connect to TIPC topology server
send subscription for {0,1,2^31-1}, time limit = none
send subscription for {1,0,0}, time limit = 60 seconds
set set of available nodes to "empty"
loop
    receive event
    if (event == publish)
        add node to set of available nodes
    else if (event == withdraw)
        remove node from set of available nodes
    else
        send subscription for {1,0,0}, time limit = 60 seconds
        print list of available nodes
end loop
```

The use of a dummy subscription having the name sequence `{1,0,0}` allows the

application to print the desired information at required time without having to use additional timers or threads of control, and without having to deal with the complexities of determining how long to continue waiting when the main thread is awoken to process a publish or withdraw event. The name sequence specified by the dummy subscription can be anything that the application designer knows will not generate any publish or withdraw events, and need not be {1,0,0}.

Note: This code ignores processing overhead that will result in each successive display occurring slightly more than 60 seconds after the previous one. This could be compensated for by measuring the overhead and subtracting it from the specified time limit each time the dummy subscription is re-issued.

[END OF DOCUMENT]